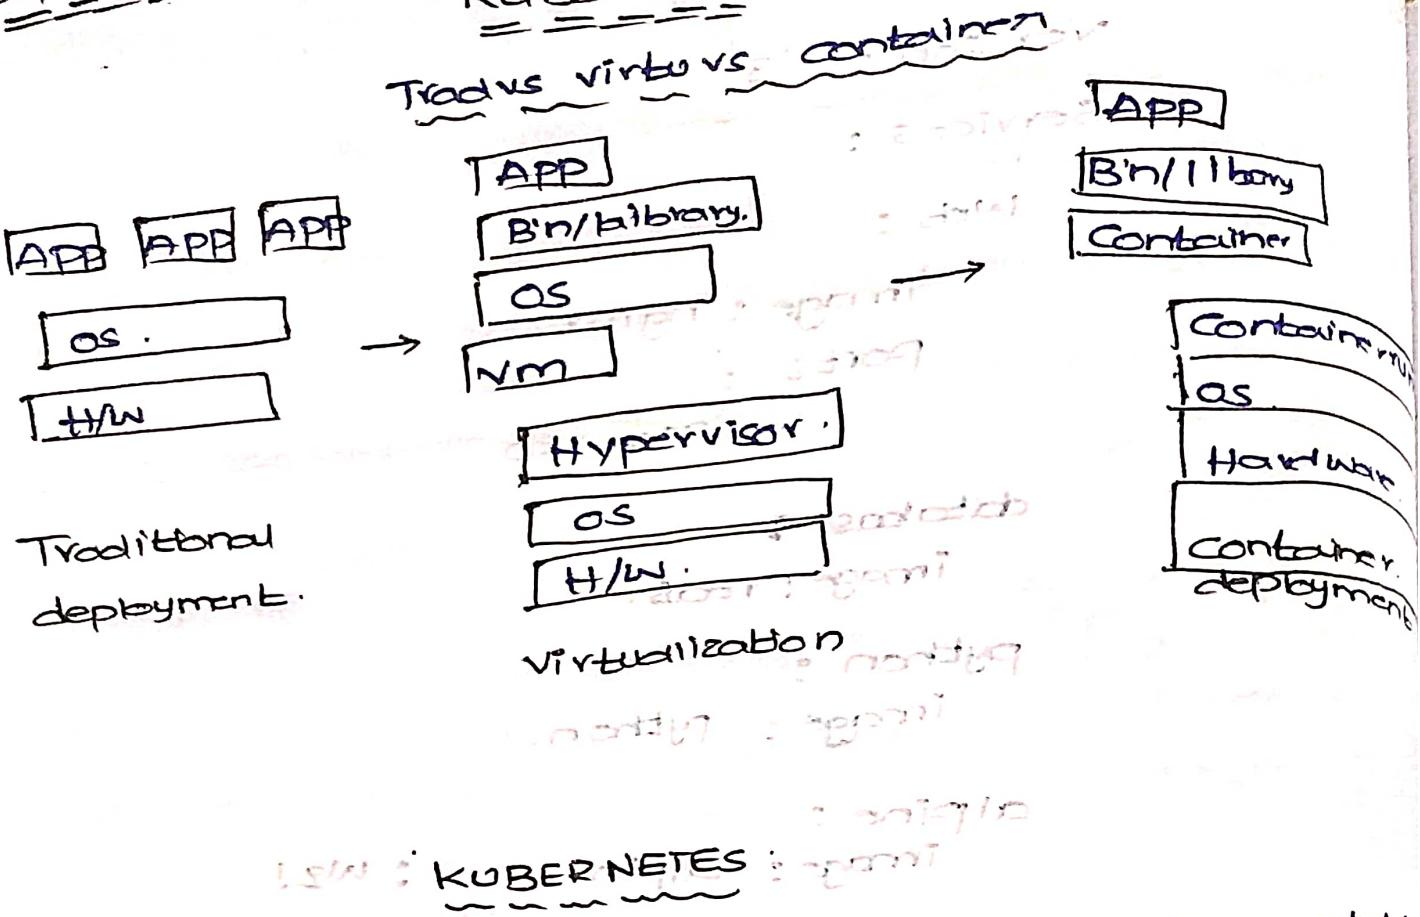


30/12/22

Kubernetes



→ Kubernetes is a portable, extensible, open source platform for managing containerized workloads and services that facilitates both declarative configuration & automation.

- Kubernetes also known as - K8S (or) "kube"
- Kubernetes cluster can span hosts across on-prem public, private or hybrid clouds

Features of Kubernetes:

→ Automated rollouts and rollbacks.

Service discovery and load balancing.

Storage orchestration.

Secret and configuration management

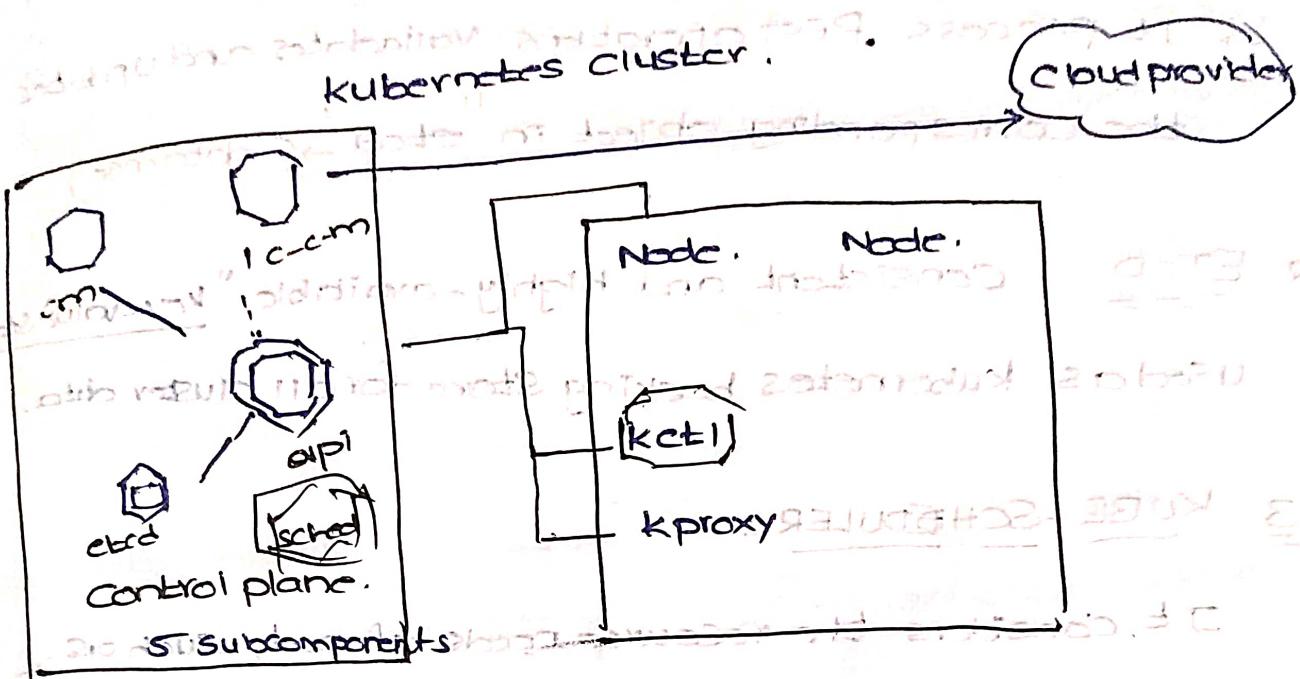
↳ possible outcome

Automatic bin packing

IPv4 & IPv6 dual stack

Self-healing [scale up and down]

Architecture & Components



→ Node → VM → Container (Nested)

* A Kubernetes cluster consists of a set of worker machines called "nodes", that run containerized application.

→ The worker node(s) host the pods that are the components of the application workloads.

→ The control plane manages the worker nodes and pods in the cluster.

→ Cluster usually runs multiple nodes; providing

"Fault-tolerance & High Availability"

* Control plane Components:-

1. KUBE-APISERVER :- IT IS THE MAIN MANAGEMENT OF THE ENTIRE CLUSTER, HANDLING INTERNAL AND EXTERNAL REQUESTS.

→ IT PROCESS REST OPERATIONS, VALIDATES AND UPDATES THE CORRESPONDING OBJECT IN "ETCD" → (DATABASE).

2. ETCD : CONSISTENT AND HIGHLY AVAILABLE "KEY-VALUE" USED AS KUBERNETES BACKING STORE FOR ALL CLUSTER DATA.

3. KUBE-SCHEDULER

IT CONSIDERS THE RESOURCE NEEDS OF POD, SUCH AS, CPU (OR) MEMORY ALONG WITH THE HEALTH OF CLUSTER.

→ IT SCHEDULES THE POD TO AN APPROPRIATE COMPUTER.

4. KUBE-CONTROLLER-MANAGER:-

IT RUNS THE CONTROL PROCESS. A CONTROL PROCESS IS A LOOP

THAT FOCUS ON MAKING THE DESIRED STATE EQUAL TO THE CURRENT STATE FOR ANY APPLICATION IN ANY GIVEN INSTANCE.

5. CLOUD-CONTROLLER-MANAGER:-

IT LINKS YOUR CLUSTER IN TO YOUR CLOUD PROVIDER.

API and separates out the components that interact with that cloud platform from components that only interact with your cluster.

* Node Components

1. KUBELET :- An agent that runs on each node in the cluster. It makes sure that containers are running in a pod.

→ When the control plane needs something to happen in a node, the kubelet executes the action.

2. KUBE-PROXY :

→ It is a network proxy that runs on each node in your cluster.

→ It maintains network rules on nodes. These rules allow network communication to your pods of your cluster.

3. Container Runtime:

→ It is a software that is responsible for running containers.

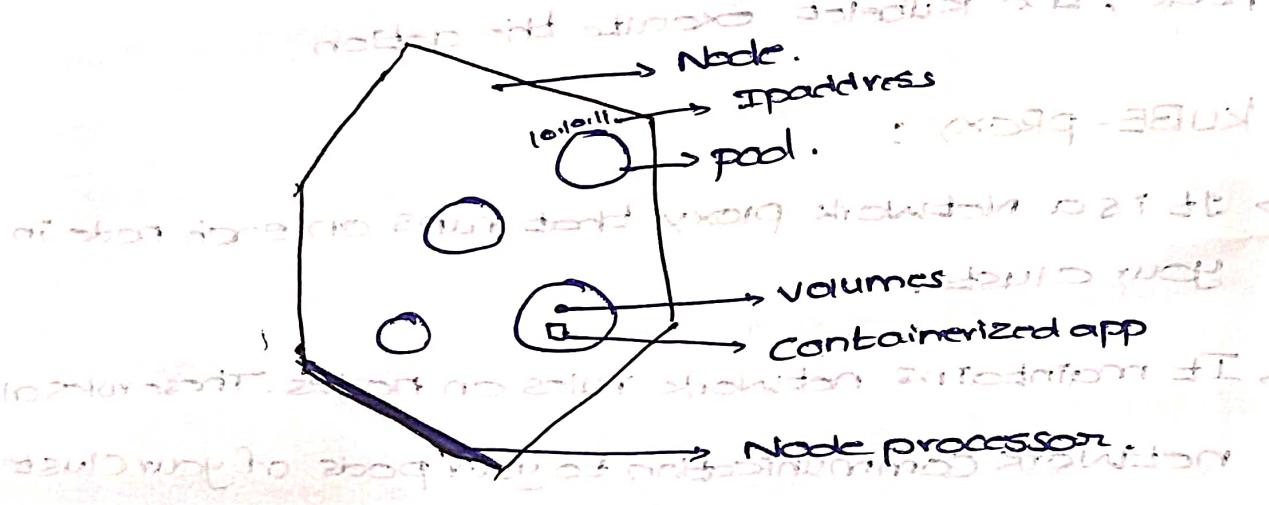
→ Kubernetes supports Docker, Containerd, CRI-O,

and Kubernetes' CRI (Container Runtime Interface).

Note: Every "node" we have to install docker on it.

Administrator can configure different docker configurations at the beginning.

- Kubernetes Nodes & pods :-
- A node is a worker machine in K8s and may be (vm or) physical machine, depending on the cluster.
 - K8s runs on your workload by placing containers in pods to run on Nodes. A Node can have multiple pods running on it.
 - A Node includes Kubernetes, a container runtime.
 - Node names are uniqueness in the cluster.



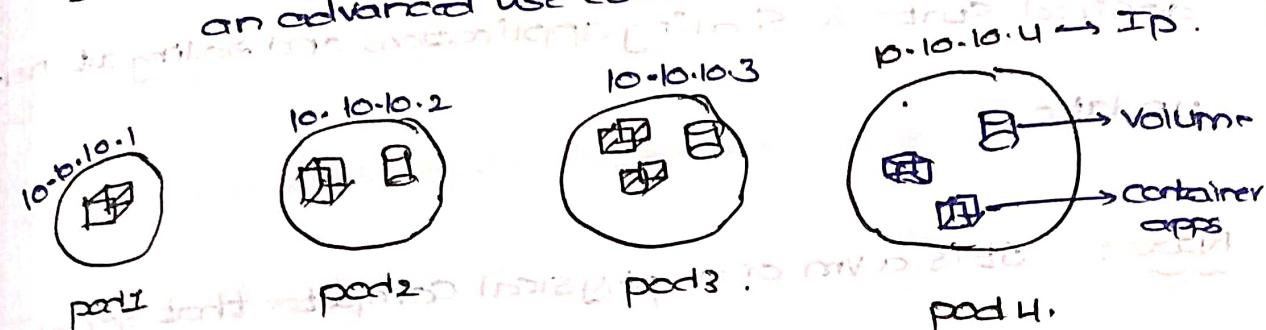
PODS

- A pod is the smallest & simplest K8's object. It represents a single instance of a running process in a cluster.
- pods contain one (or) more containers, such as Docker containers.

Docker containers

- pods run multiple containers; the containers are managed as a single entity & share the pod resources.

- pods also contain shared networking and storage resources for their container.
- running multiple containers in a single pod is an advanced use case.



2/1/23: without installing k8s, easiest way to start k8s is

MINIKUBE

- minikube is a utility you can use to run kubernetes

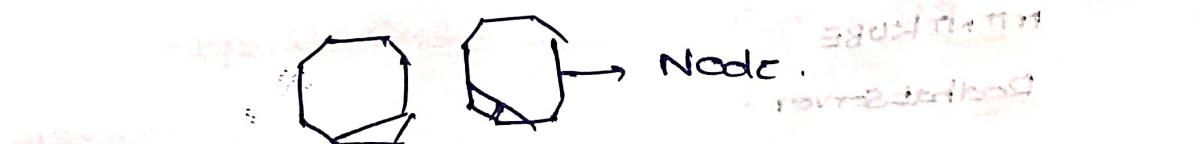
(k8s) on your local machine.

- It is a light weight kubernetes implementation that creates a vm on your local machine and deploys a simple cluster containing only one node.

- minikube CLI provides basic bootstrapping operations

for your cluster, including start, stop, status & delete.

minikube Architecture



Container

Network Processor

R-cluster

Control Plane :- It is responsible for managing the system resources and maintaining application. Such as scheduling application, maintaining application.

→ The control plane coordinates all activities in your cluster. Such as scheduling application, maintaining application, desired state & scaling applications and rolling out updates.

Node :- It is a VM or a physical computer that serves as a worker machine in a Kubernetes cluster.

→ Each node has a "kublet and container runtime" for managing the node and handling container applications.

* minikube installation:

1. Installation of minikube [Windows, Mac, Linux] -> 4GB RAM
PRE- Requirements :- 1. 2CPU or more.

2. 2GB of free memory

3. 20GB of free disk space

4. Container (or) VM manager (Docker, VirtualBox, VMWare, KVM, Hyper-V)

* Go to AWS instance (EC2) → Launch instances

MINIKUBE

RedhatServer

t2-medium [per hour → 4 rupees]

Create SG :-

1x20 GB → [Launch Instances]

→ Select Instance → Security → Click on Security Group
Edit Inbound rules → Add rule → open all TCP, anywhere.
Save

2. Connect Instance → Open hosts file

→ To change hostname: hostname minikube.

→ vi /etc/hostname [to update hostname]
→ keep Computer name (minikube) always!

exit. → for update of computer name.

→ sudo -i

→ yum update -y [Install updates]

→ cat /etc/redhat-release [Server version]

→ ifconfig [inet: 172.31.56.87]

→ yum install net-tools vim -y
for ifconfig command.

→ ifconfig [inet: 172.31.56.87]

→ vim /etc/hosts [mapping ipaddress to hosts]

IPaddress

name → IPaddresses
IPaddresses → name } mapping

Server hostname

IPaddress + server

→ systemctl stop firewalld [to stoping the firewalls]

→ getenforce [to check status of firewall]

* Docker Installation:

= yum install -y yum-utils.

1. yum install -y yum-utils.

2. cd /etc/yum.repos.d/

vim docker-ce.repo [created repository]

↓
add six lines of configuration

about [minikube] plugin by.

3. yum install docker

4. systemctl start docker
enable --> start

5. Go to google → minikube install (minikube stable)

Select Installation Step [Linux] Stable RPM package

→ minikube download [stable minikube]
→ exit

→ curl -lo https://storage.googleapis.com/minikube/stable/minikube

→ Run the file: sudo rpm -Uvh minikube-latest

2. Start your cluster: → docker driver

→ sudo + command

→ minikube start --driver=docker --force

→ minikube

→ minikube status.

Creating docker
container

`sudo minikube kubectl -- get po -A` [Running]

Manage your cluster:

- * → `Sudo minikube pause`
- `Sudo minikube unpause`
- Stop the minikube:- `sudo minikube stop` " " `start`

Additional notes (E) minikube vs Docker

3/1/23

KUBEADM Installation

KubeADM: Kubeadm is a tool used to build multi-node kubernetes (k8s) cluster.

→ Kubeadm provides kubeadm init and kubeadm joins as best-practices "fast paths" for creating kubernetes cluster.

[connecting] → kubeadm provides the actions necessary to get a minimum viable cluster up & running.

prerequisites

→ A compatible Linux host (Debian (or) Redhat).

server → master node should have 2GB RAM & 2CPU.

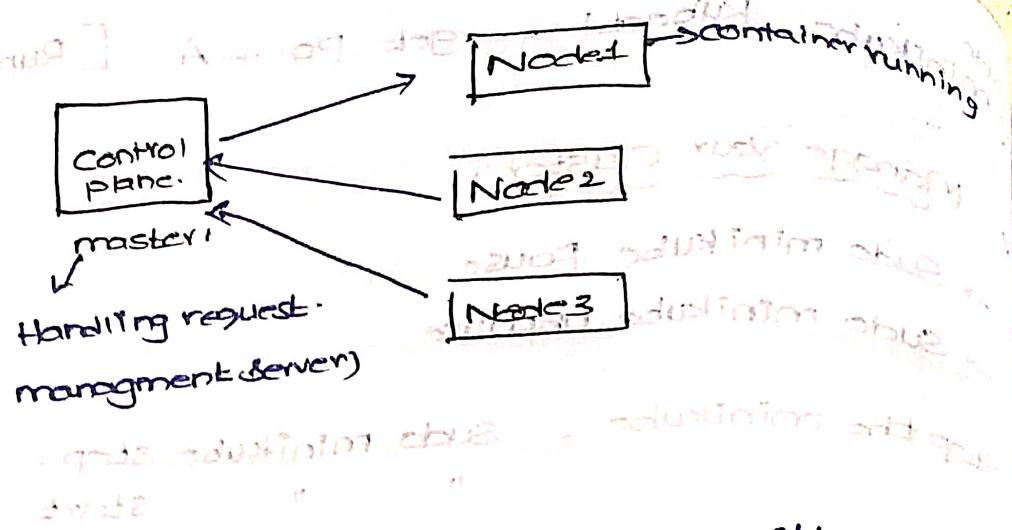
client → worker node " " 1GB RAM & 1CPU

→ Full N/W connectivity b/w all machines in the cluster

→ Unique "hostname", macaddress and product-id for every node.

→ certain ports are open to your machine.

→ Must disable Swap(virtual memory)



* Step1: Creating EC2 instance (3) with one master and 2 nodes.

→ Launch instance, the Nb of instance "3", "Linux"

1 instance → Linux, free tier.

N/W setting edit → SG → Kubernetes

edit inbound rules: 1. Allow all traffic on port 22.

→ Open All TCP & anywhere address.

edit inbound ports: after editing port 2207, adding 2208 to 2210

→ Storage 20GiB. → Launch instances.

1. master 2. Node1 3. Node2 [3 instances]

→ Step2: Stop the master machine.

Action → Instances setting → changing instance types. →

(t2-medium) → Apply. → Start instances.

Note:

Master Configuration: Connect master instance →
putty → Login → add, edit, remove configurations

Step 1: change hostname → hostname master.
type hostname in a file.
vi /etc/hostname [master] : wq!

exit.

Sudo -i

Step 2: yum update -y [Server update]

Step 3: yum install net-tools vim -y

ifconfig

Step 4: vim /etc/hosts

172.31.80.129 master. : wq!

Step 5: Installing Docker

1. Vim /etc/yum.repos.d/docker-ce.repo

↓

2. add six lines. [: wq!]

3. yum install -y plugin -y

→ Systemctl start docker

" enable "
" Stop "

→ Steps: Go to google → kubeadm install.

Installing kubeadm, kubelet and kubectl

Redhat based distribution [commands]

→ copy all commands, → systemctl enable kubelet

→ Step 1: Creating a cluster on master. → ifconfig
if eth0

[root@master] kubeadm init --pod-network-cidr
10.16.0.0/16 --apiserver-advertise-address =
31.80.129.7 [master Node IP address] Error resolv

→ root user → export KUBECONFIG = /etc/kubernetes/admin

→ Token → Joins command copy.

→ How do you know nodes join in master.

kubectl get nodes

→ Installing pod N/W

kubectl apply -f https://github.com/coreos/flannel/raw/master/Documentation/kube-flannel.yaml

port issue:-

kubectl get nodes

Node configuration

Connect Node1 Instance [Ip+key] → Login. [both]

Nodes Should be connected parallelly. [2 terminals]

Step1. yum update -y

Step2. yum install net-tools vim -y

Step3. host name of Nodes

vi /etc/hostname [Node1] : wq!

ifconfig.

vi /etc/hosts.

172-31-81-181 Node1

172-31-87-254 Node2 : wq!

Step4. ~~cat~~ cat /etc/hosts.

Adding Node machine address in to master file
[root@master]

master machine: 172-31-80-129 master
node1 node2

Steps: Install docker on Node1 Node2.

1. vi /etc/yum-repos.d/docker-ce.repo

↓

2. addsixline

3. yum install --- plugin ---y

→ Systemctl start docker | bothNodes
" " enable "

Steps:

*. Snapshot (-a name) [stop VM]

→ Google → Redhat distribution → commands. [400+]

→ Delete a file → rm /etc/containerd/config.toml

Kill and Join Commands! [copy]

kill [pid]

join [pid].run

exit [pid].run

exec [pid].run

lsblk

lsblk -f -E -SFI

fdisk

fdisk -l -SFI

fdisk -l -SFI

Errors on master cluster Node

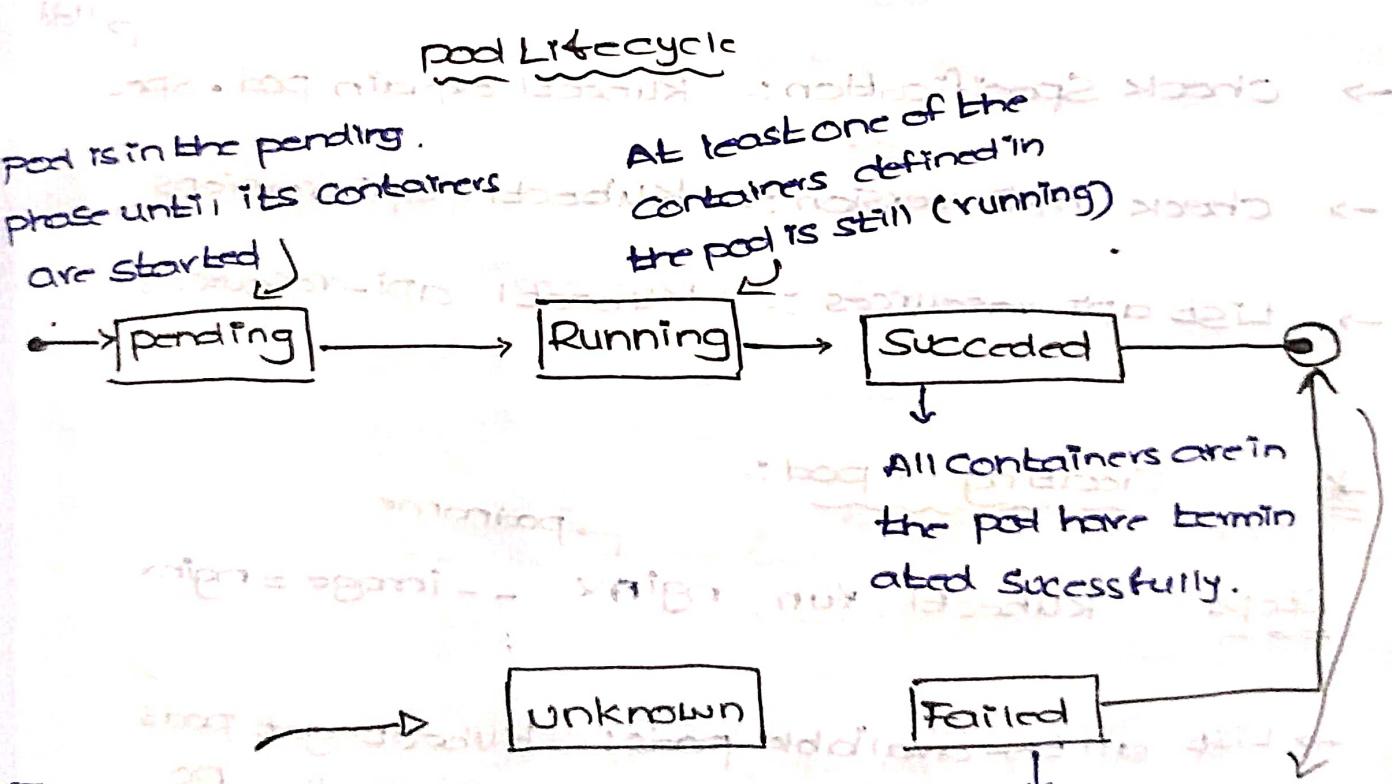
- kubeadm reset · | ~~Container runtime not run~~
1. rm /etc/containerd/config.toml [on master]
 2. Systemctl restart containerd

4/1/23

PODS :: A pod is the smallest & simplest Kubernetes object.

- pods represent a "single instance of a running process in your cluster".
- pods contain one or more containers, such as Docker and standard Linux containers.
- When a pod runs multiple containers, the containers are managed as single entity & Share the pod's resource.

Note: Running multiple containers in a single pod is an advanced usecase.



The state of the pod is shown as unknown when the Kubernetes API server stops reporting to the API server (management Server). The pod had terminated unsuccessfully.

* Creating a pod :-

- kubectl get nodes.
- Complete details of node :- kubectl get nodes -o yaml
- 3-service need to running state. While working with cluster, we can check status of pods with `get pods`.
Sudo `systemctl restart containerd`
" " " " restart kubelet
- How to check the resources ?

Syntax :- `kubectl explain resources {pod, N/W, volume}`

`kubectl explain pod`

- Check Specification :- `kubectl explain pod.spec`
- Check API-Version :- `kubectl api-versions`
- List API Resources :- `kubectl api-resources`

* Creating a pod :-

Step1: `kubectl run nginx --image=nginx`

- List all the available pods : `kubectl get pods`
- Check complete pod details :- `kubectl describe pods nginx`
1. `kubectl describe pods nginx`
2. `kubectl get pods -o yaml`

→ Create a pod with the help of yaml file

Create a pod with yaml file.

Step1

3-services → Running.

pods write in microsoft visual codes [yaml]

base of k8s

Step2 install microsoft visual studio → VS download
free,

Newfile:-

apiVersion: v1

kind: pod

metadata:

name: my-pod

labels:

app: myapp

Spec:

containers:

-name: myapp-containr

image: nginx

command: ['sh', '-c', 'echo Hello Ragu...
sleep 3600']

→ Create a file → vim pod.yaml → Same code

kubectl create -f pod.yaml

kubectl get pods.

kubectl get describe pods my-pod

* Creating multiple containers in one pod
 → vi mpods.yaml

```

apiVersion: v1
kind: pod
metadata:
  name: my-site
  labels:
    app: web
spec:
  containers:
    - name: front-end
      image: nginx
      ports:
        - containerPort: 80
        - name: rss-reader
          image: nickchar/rss:php-nginx:v1
          ports:
            - containerPort: 88
    : w2!
  
```

→ `kubectl create -f mpods.yaml`
`kubectl get pods`

⚡ Delete a pod
`kubectl delete pods my-pod`

Labels: ~~are written in the underline~~ metadata resources

- Labels are key-value pairs that are attached to objects, such as pods.
- It is used to specify identifying attributes of objects that are meaningful and relevant.
- Label can be used to organize and to select subset of objects.

Ex: 1000 pods

"metadata": {

 "key1": "value1"

 400 - pro-team

 "key2": "value2"

 200 - dev - team

 "labels": {

 "key1": "value1"

 }

 "key1": "value1",

 "key2": "value2"

 "key2": "value2"

 "key3": "value3"

} }

 "key3": "value3"

 "key4": "value4"

* Annotations:

It is also "key-value" pair for connecting non-identifying metadata with object.

These are not used to identify and select object.

metadata": {

 "annotations": {

 "key1": "value1"

 "key2": "value2"

}

* Example of label-yml

→ vim label-yml

apiVersion: v1

kind: pod

metadata:

name: label-demo

labels:

environment: production

app: nginx

ver: 1.4

proj: airtel

} : "label-demo"
} : "airtel"

Spec:

containers:

-name: nginx

image: nginx:1.14.2

ports:

containerport: 80

: W2!

→ kubectl create -f label-yml

→ kubectl describe pods ~~label-demo~~ label-demo

5/1/23 To Connect a pod
→ podname
Kubectl exec -it my-pod -- bash
[my/pod] --exit
→ Label Selectors:

- Labels doesn't provide uniqueness, In general, we expect many objects to carry the same labels.
- The label selection is the core grouping primitive in kubernetes.
- It support two types:

- Equality Based Requirement:
- Equality (or) Inequality-based requirements allows filtering by label keys & values.
 - matching objects should satisfy all the specified labels
 - Supported operators are: $=$, \neq , $!=$.

- Set-Based Requirement:
- It allow filtering keys according to a set of values.
 - Supported operators are: in , $notin$, $exists$.

* How to filter pods with labels:-

Kubectl get pods.

kubectl get pods --show-labels

→ to list labels for filtering

EOR

Kubectl get pods -l environment=production

SBR :- kubectl get pods -l environment in (production)

↳ environment in production

→ kubectl get pods -l 'environment in (production,

development)'

* Annotation Example:

vim /annotation.yaml

apiVersion: v1

kind : pod

metadata :

name : annotation-demo

annotation :

provider : raju

mailId : rrajuuu4@gmail.com

imageRegistry : "https://hub.docker.com/"

* Spec :-

Containers :

-name: nginx

-image: nginx : 1.14.2

ports :

-containerPort : 80 → 80



Example

→ Given object - metacalls → label
↓
proj - Airtel

region - produ container details.

web - Apache

Annotation: URL from where WebServer download.

kubectl create -f annotation.yaml

kubectl get nodes

kubectl describe pods annotations-demo

detail → job run on Node decided by master node.

action step 1: node 2

NAMESpaces : [Important]

→ Namespaces provide a mechanism for isolating group of resources within a single cluster.

[or]

→ It is a way to organize cluster into virtual sub-clusters. They can be helpful when different teams or projects share a Kubernetes cluster.

→ Names of resources need to be unique within a namespace, but not across namespaces.

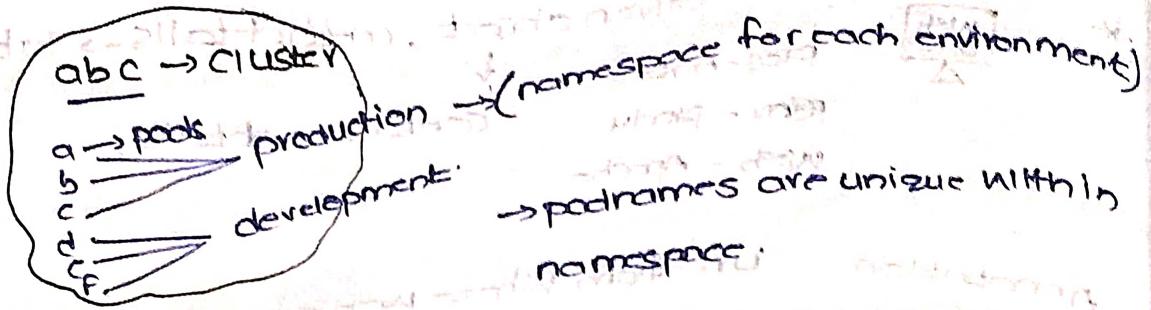
→ Kubernetes starts with four initial namespaces.

→ default

→ kube-node-local

→ kube-public

→ kube-system



→ `kubectl get ns` [namespaces] → To list namespaces

→ How to create a namespace :-
→ Namespace → Nameofns.

→ `kubectl create namespace development`

`kubectl get nodes`.

[nodes] → nodes available

→ To switch namespace:-

→ `kubectl config set-context --current --namespace=development`

= development.

→ `kubectl get pods`. [No resource (or) pods not found.

→ Create a pod under development :-

→ `vim pod.yaml`

→ `kubectl create -f pod.yaml`

→ `kubectl get nodes`.

Switch back:

→ `kubectl config set-context --current --namespace=default`

→ switch to development:

→ Create a pod on development namespace:

* vim namespace.yaml

apiVersion: apps/v1

kind: Deployment

apiVersion: v1

kind: pod

metadata:

name: label-demo.

namespace: development

labels:

environment: production

↓ important: app: nginx, name of pod, etc

proj: airtel

Spec:

containers:

-name: nginx

image: nginx:1.14.2

ports:

containerPort: 80

→ kubectl create -f label.yaml

kubectl get nodes

Switch to kubectl config set-context --current
--from=cluster of home@airtel and quboot

--namespace=default

→ kubectl get all -A

- * To delete a namespace:-
 - kubectl get ns
 - kubectl delete ns development
 - kubectl delete pods → kubectl get pods
 - kubectl delete pods annotations-demo [label demo]
 - my-pod [deleting multiple pods]
 - kubectl get nodes

5/1/23.

- Kubernetes Deployment → main application H-V
 - Deployment → maintain application
 - Deployment → maintain application
- JE is used to create (or) modify instance of the pods, that hold a containerized application.

- A deployment provides declarative updates for "pods and Replicaset".

Usecase of deployment :-

- Create a deployment to rollout a Replicaset.
- Declare the new state of the pods.
- Rollback to an earlier deployment reversion.
- Scaleup the deployment to facilitate more load
- pause the rollout of a deployment.
- Use the status of the deployment.
- Clean up older replicsets.

REPLICASETS :- (RS)

- A Replicaset (RS) is a Kubernetes object used to maintain a stable set of replicated pods running within a cluster at any given time.
- It is often used to guarantee the availability of a specified number of identical pods.

Example :- vim deployments.yaml

apiVersion : apps/v1

kind : Deployment

metadata:
name : nginx-deployment

labels:

app : nginx

Spec :

replicas : 2

selector:

matchLabels:

app : nginx

Template:

(with the same metadata, labels, spec)

labels:

app : nginx

Spec :

ports:

- containerPort : 80

Container :

- name : nginx

: 112 !

Image : nginx : 1.14.2

→ `kubectl create -f deployments.yaml` → It gives pod only.

→ `kubectl get pods`

→ Connect a pod :-

→ `kubectl exec -it -n nginx deployment - bash`

→ exit.

* To check the roll out status :-

→ `kubectl get deployments` → Name
→ `kubectl rollout status deployment/nginx-deployment` → mont.

→ `kubectl get rs`

→ Label name : `kubectl get pods --show-labels`

* To modify the Replicaset value :

Goto → vi `deployments.yaml` → under Speciation change

the value "replicas: 5"

Note : `kubectl apply -f deployments.yaml`

↓

only warning message.

→ `kubectl get deployments`

→ `kubectl get pods` → (5 pods displays)

* How to rollout to back (5-2) :-

`kubectl get rs`

→ To rollout replicasets :- [yaml fixed]

kubectl rollout undo deployment/nginx-deployment

* Rolling back a deployment

- Sometimes you may want to roll back a deployment

For Example : When the deployment is not stable, such as

"Cross-looping"

→ To check history of deployment :-

kubectl rollout history deployment/nginx-deployment

↓
Revision = 1 [1 time rollout]

→ kubectl rollout history deployment/nginx-deployment

-- revision = 1

→ Rollback to a specific revision :- [-to]

kubectl rollout undo deployment/nginx-deployment

↓
-- to-revision = 1

→ To verify deployment description :-

kubectl describe deployment/nginx-deployment

* Scaling a Deployment :- [Scale]

→ Kubernetes Scale deployment/nginx-deployment

-- replicas = 10

→ kubectl get deployments
→ (Only update, set limitation CPU limits)

* Pausing & Resuming a deployment

→ kubectl rollout pause deployment/nginx-deployment

kubectl get deployments

→ Change the Image of the deployment [setImage]

→ kubectl set image deployment/nginx-deployment

nginx = nginx : 1.16.1

→ kubectl rollout history deployment/nginx-deployment

→ kubectl get rs.

→ kubectl describe deployment nginx-deployment.

Traditional method

Dep → Web 2.2

↓
2.4.

Command → yum update -y

Web: 2.4

Increase, (CPU (or) memory)

takes time

Kubernetes

SetImage = 2.4,,

--cpu = 8,

--memory = 512

Kubectl Set resources deployment

Resource:

* `kubectl set resources deployment nginx -c=nginx --limits="cpu=200m, memory=512Mi"`

nginx-deployment

Kubectl describe deployment nginx-deployment

→ kubectl get rollout status deployment/nginx-deployment

→ Rollback the deployment rollout and observe a new

replica set coming up with all the new updates.

→ kubectl rollout resume deployment/nginx-deployment

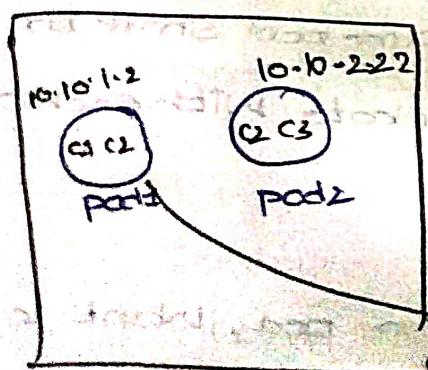
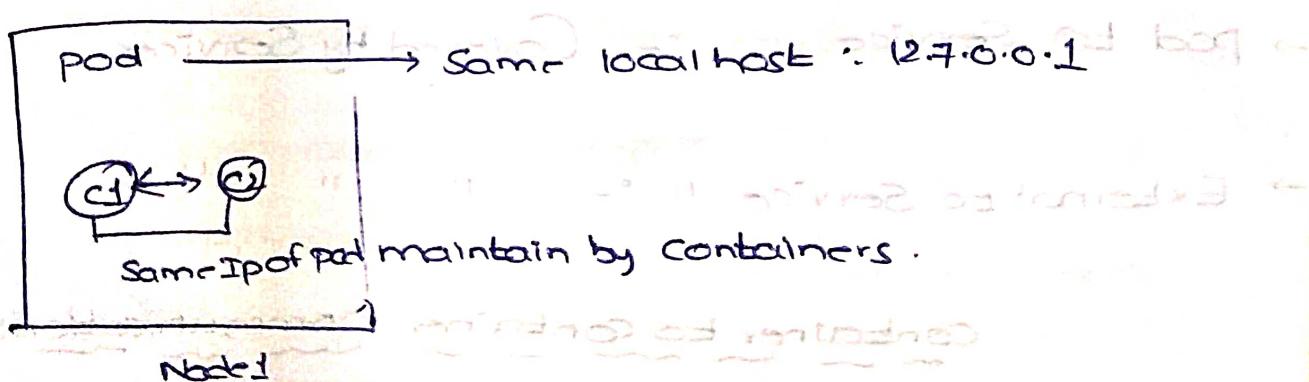
→ kubectl get rs -l app=nginx

→ kubectl get rs

→ kubectl rollout status deployment/nginx-deployment

→ kubectl rollout status deployment/nginx-deployment

7/1/23: KUBERNETS NETWORKING [cluster N/W]



→ Every pod maintains different IP

Port1 → pod2 → communication

Node1 → Node2 → communication

Overlay N/W

external use (Service) exposure
→ (User) outside of N/W

K8's Networking

- It allows Kubernetes components to communicate with each other and with other applications.
- K8s Networking is based on a "Flat Network" structure that eliminates the need to map host ports to container ports.
- 2 types of communication:-
 1. Inter-node communication.
 2. Intra-node communication.

4 Networking problem To Address:-

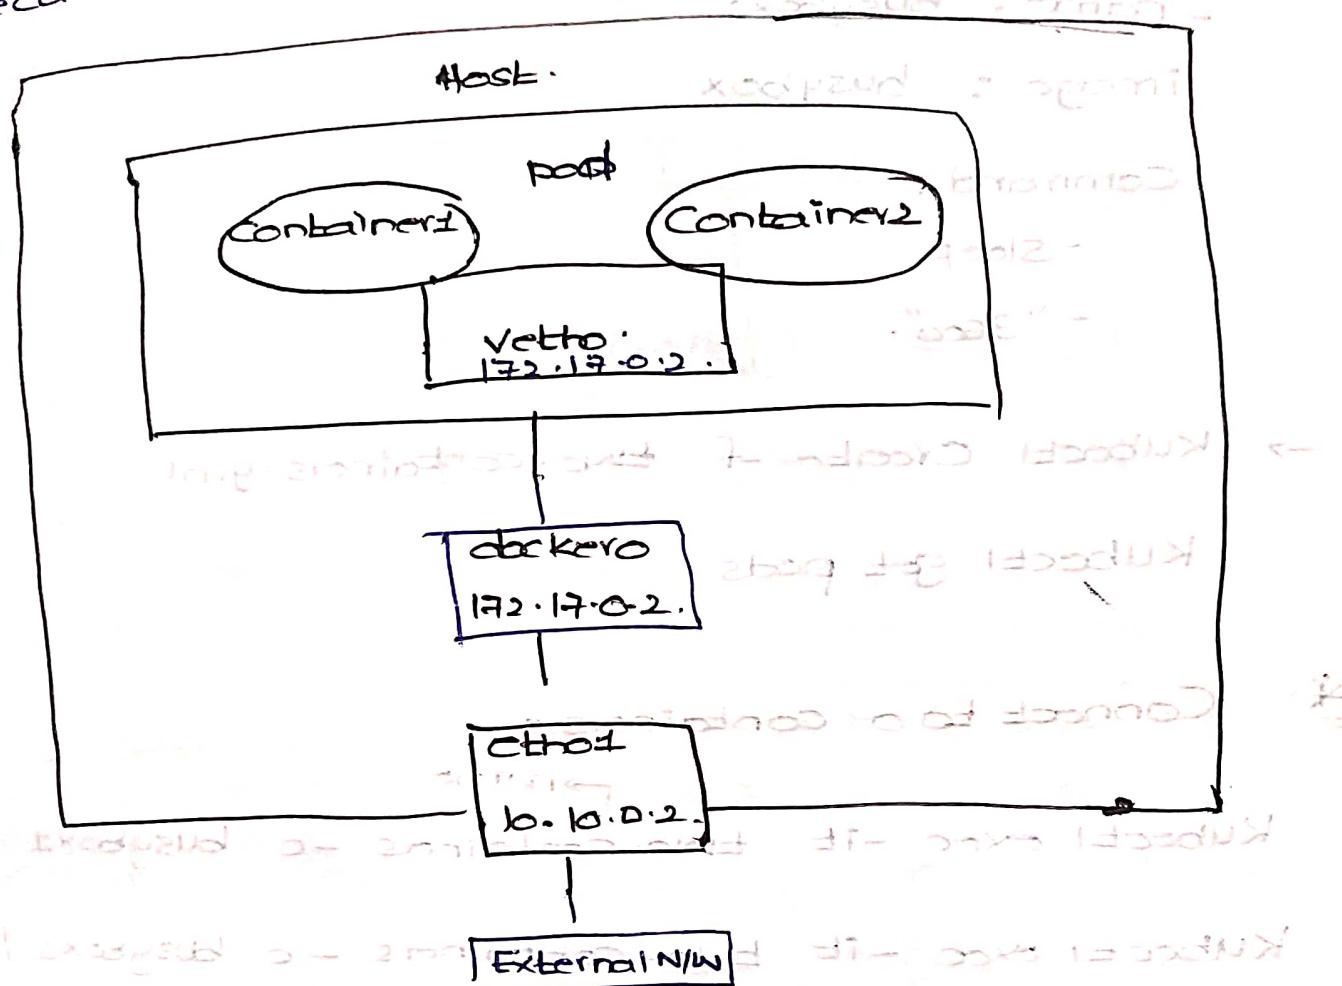
- Container to Container communication: Solved by pods and local host cmu.
- pod to pod :- The primary focus of this document.
- pod to Service :- Covered by Services.
- External to Service :-

Container to Container communication:-

- multiple containers in the same pod share the same IP address. They can communicate with each other by addressing local host.

For Example, if a container in a pod, want to reach

another container in the same pod on port 8080, it can use the address localhost:8080. You couldn't have two containers in the same pod which expose port 8080, because there would be a conflict.



* Create a pod with two containers?

vi two-containers.yaml

apiVersion: v1

kind: pod

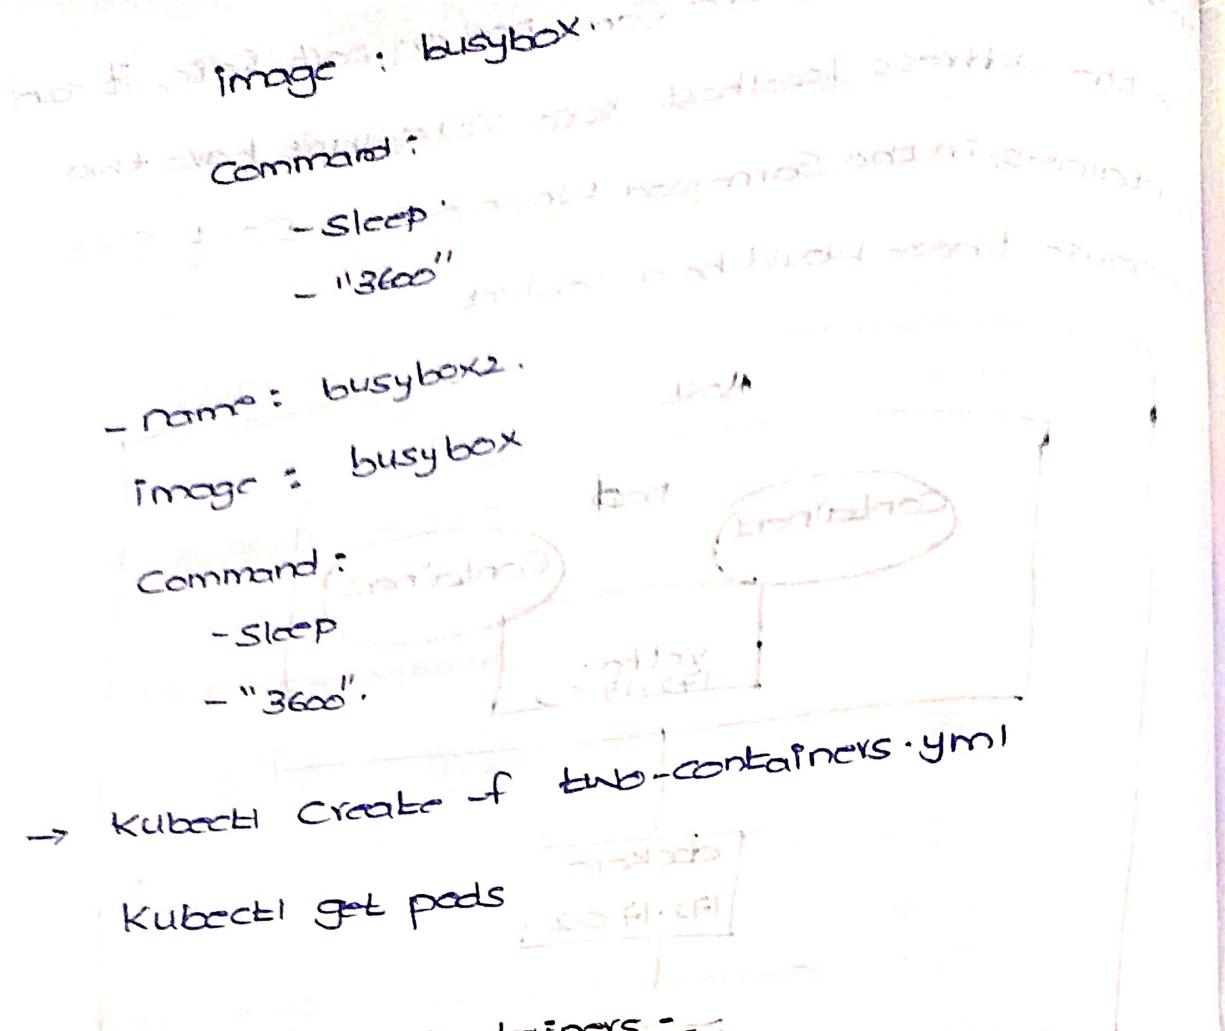
metadata:

name: two-containers

Spec:

Containers:

- name: busybox1



1. Connect to a container:

Kubectl exec -it two-containers -c busybox1 /bin/sh

Kubectl exec -it two-containers -c busybox2 /bin/sh

↓
 Result verify IP addresses are same
 (Container 1 and Container 2 have same IP address)

2. Pod to pod communication:

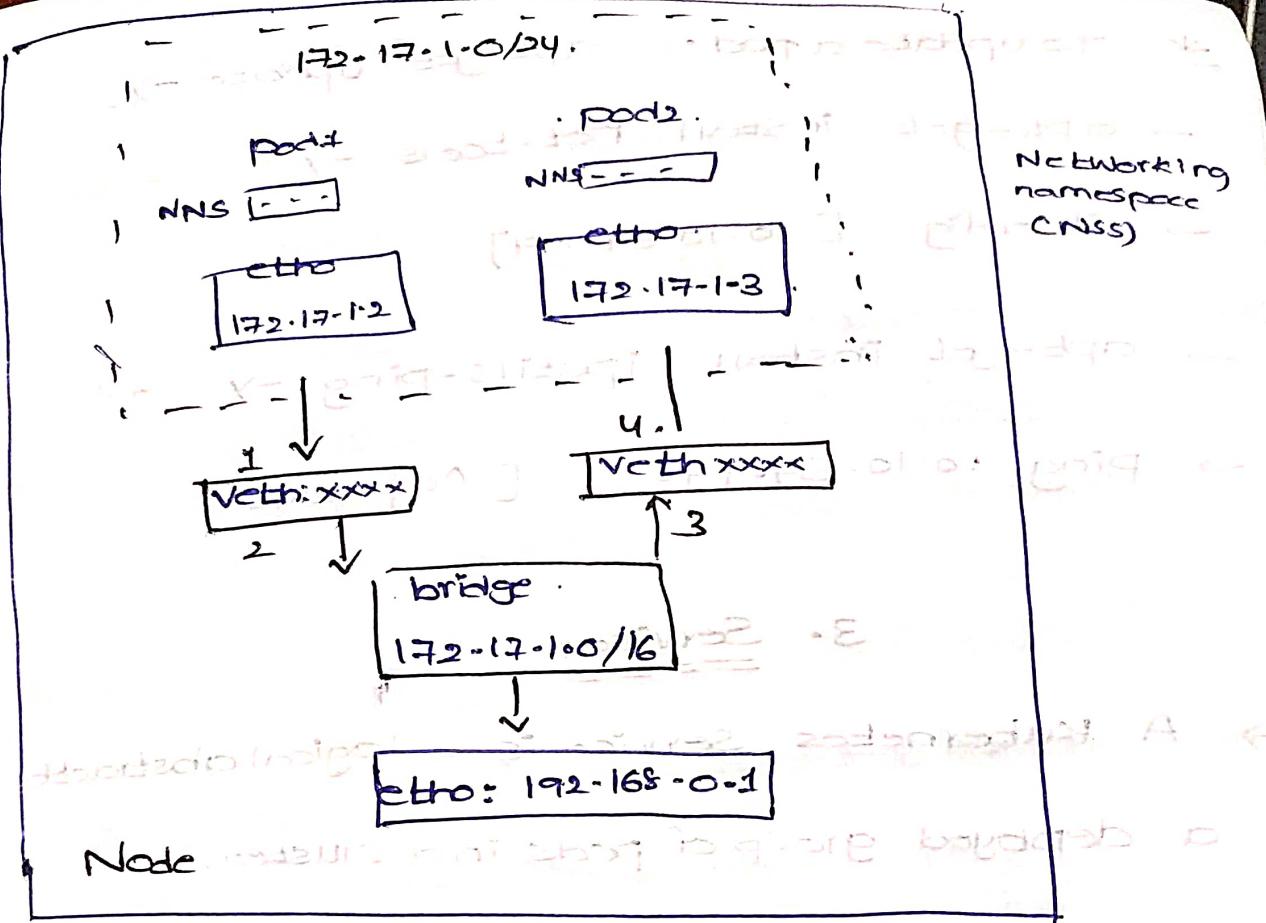
→ In Kubernetes, each pod has its own IP address.

At a very primitive level, pods can communicate with each other using their IP address.

→ Pod to pod communication requirements:

Layer 2 (Switching) Solution.

Layer 3 (Routing), overlay solution.



* Create two pods?

= To run two different nginx version on both nodes

→ Create a pod 1 : `kubectl run nginx1 --image=nginx`

2. `kubectl run nginx2 --image=nginx`.

→ `kubectl get pods`.

→ To check IP address of pod: `get pods -n kube-system`

`kubectl describe nginx1 | grep -i ip` [IP = 10.10.69.197]

`kubectl describe nginx2 | grep -i ip` [10.10.69.198]

→ Connect to pod1: `exec -it nginx1 -- bash`

`kubectl exec -it nginx1 -- bash`

* To update a pod :- `apt-get update -y`

→ `apt-get install net-tools -y`

→ `ifconfig [10.10.69.197]`

→ `apt-get install iputils-ping -y`

→ `ping 10.10.69.198 [^c] exit`

3. Service :-

→ A Kubernetes Service is a logical abstraction for a deployed group of pods in a cluster.

→ It is exposing an application running on a set of pods as a Network Service. Which can be reached by a single fixed DNS name or IP address.

→ The set of pods targeted by a Service is usually determined by a "Selector".

→ There are various Service types options for exposing a service outside of your Cluster IP address.

1. ClusterIP : Expose a service which is only accessible from "within the cluster"

2. NodePort : Expose a service via a static port on each node's IP

Net port must be in the range of 30,000 - 32,767

Load Balancer: Exposes the service via the cloud provider.
Load balancer.

External Name: maps a service to a predefined external.
Name field by returning a value for the Cname record
↓
(Canonical)

④ 8/123

ClusterIP:

It exposes a service which is only accessible from within
in the cluster.

Ex:- Create a deployment `clusterdeploy.yaml`

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

Spec:

replicas: 2

Selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

Spec:

Contains:

- name : nginx,

image : nginx-latest

ports:

- containerport: 80

: 1112!

→ kubectl create -f cluster-deployment.yaml

kubectl get pods

* Create a Cluster IP for the deployment,

vim cluster-service.yaml

apiVersion: v1

kind: Service

metadata:

name: my-service

spec:

Type: ClusterIP

Selector:

app: nginx

ports:

- protocol: Tcp

port: 8080

targetport: 80

: 1112!

→ kubectl create -f cluster-service.yaml

Kubectl get service

Kubectl describe service my-service

→ Nodeport :- Expose a service via a static port
on each node IP address

→ Nodeport must be in the range of "30000 to 32,767"

Create a deployment :-

vim nodeport deployment.yaml

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

labels:

app: nginx-app

spec:

replicas: 1

selector:

matchLabels:

app: nginx-app

template:

metadata:

labels:

app: nginx-app

spec:

containers:

- name: nginx-np-Container

image: nginx-latest

container:

- containerPort: 80 : 1112!

* kubectl create -f nodeport.yaml

kubectl get pods.

* Create a Service:

vim nodeport-service.yaml

apiVersion: v1

kind: Service

metadata:

name: my-service-nodeport

labels:

app: nginx-app

Spec:

type: NodePort

ports:

- protocol: Tcp

port: 80

target Port: 80

node port: 31000

Selector:

app: nginx-app

→ kubectl create -f nodeport-service.yaml

kubectl get service

kubectl get describe service my-service-nodeport

* Node → IP address

* Node → IP address

→ Go to Web browser, then add Node IP address:

3600

localhost:3600

* Kubernetes Storage :-

Kubernetes volumes: Volumes in a Kubernetes is a directory which is accessible to the container in a pod.

→ Volumes allow user to store data outside of the container.

→ These are not limited to any container. It supports any or all the containers deployed inside the pod of k8s.

→ It supports different kind of storage. When in the pod can use multiple of them at the same time.

Volume types in k8s:

→ AWS Elastic Block Store

- get Persistent Disk

- azure file

- azure disk

- cephfs

- cinder

- emptydir

- hostpath

- fc (fibrechannel)

- flocker

- gitrepo

- glusterfs

- iscsi

- local

- nfs

- vsphere volume

EmptyDir :- [Temporary Directory]

→ Volume is created when a pod is assigned to a node.

It exists as long as the pod is running on that node.

Volumes is initially 'empty' and the containers in the pod can read & write the files in the 'emptydir' volumes.

→ When the pod is removed from the node, the object in the "empty Dir" is erased.

* Kubernetes explain pod-spec-volumes [command]

1. Create a pod for Empty directory.
vim emptydir-pod.yaml

apiVersion: v1

kind: Pod

metadata:

name: emptydir-one-container

Spec:

Containers:

-image: Centos:7

Command:

Sleep

"3600"

name: test-container

VolumeMounts:

mountPath:

/tmp

name: tmp-volume

VolumeMounts:

- mountPath:

[Container]

Volumes:

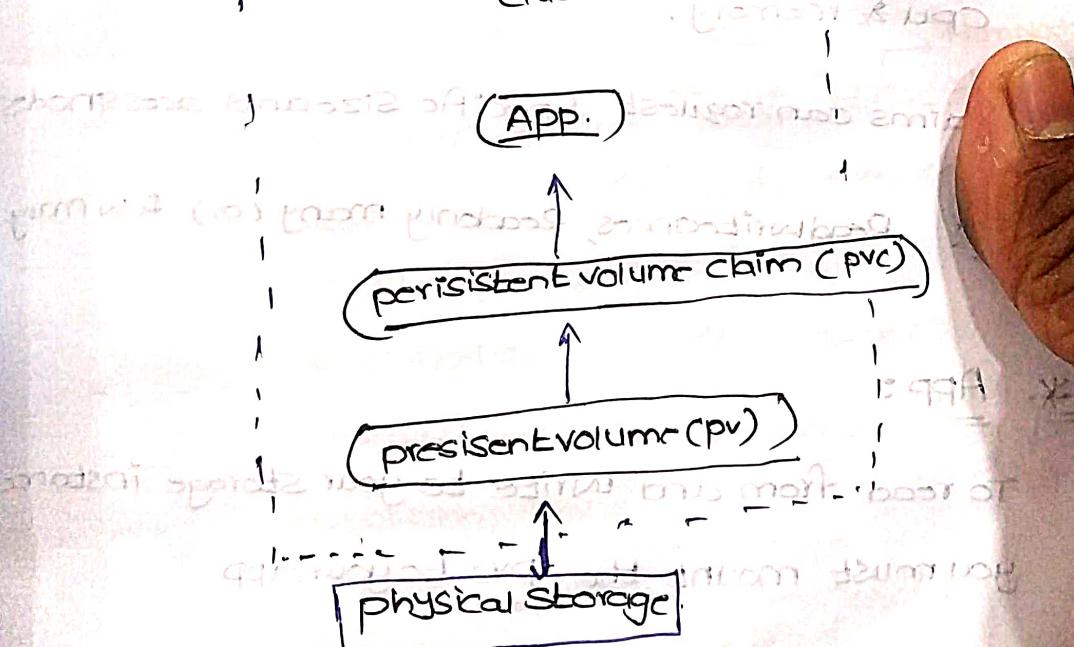
Name: tmp-volume

emptyDir: {}

Kubectl create of emptydir-pod.yaml

- Kubectl describe pod emptydir-one-container
- connect to pod: (19) and not doing anything
- Kubectl exec -it emptydir-one-container -- bash
- emptydir brew buildvorg
- Hostpath: [Important]**
- A hostpath volume mounts a file (or) directory from the hostnode's filesystem into your pod.
- The hostpath volume mounts a resource from the host node filesystem.

- The hostpath volume mounts a resource from the host node filesystem.
- pods running on the same node and using the same path in their host path volume.



* physical storage:

A physical storage instances that you can use to persist your data.

- * Persistent Volume (PV):
- PV is a piece of storage in the cluster that has been provisioned by an administrator (or) dynamically.
 - provisioned using storage class.
 - It is a resource in the cluster just like a host is and can be used to provision storage to a cluster resources.

- * Persistent Volume Claims (PVC):
- PVC is a "request" for storage by a user.
 - It is similar to a pod. pods consume nodes, resources and PVC consumes PV resources.
 - pods can request specific levels of resources like CPU & memory.
 - Claims can request specific size and access mode.

Ex: ReadWriteOnce, ReadOnlyMany (or) R/W Many.

* App:

To read from and write to your storage instance you must mount the PVC to your app.

* Cluster:

By default, every cluster is setup with a plugin to provision file storage.

- you can choose to install others add-ons, such as the one for block-storage.
- To use storage in a cluster, you must create a persistent volume claim (pvc), a persistent volume and a physical storage instance.

Note:- Different storage types have different Read-Write rules.

* Create a persistent volume (pv) :-

Step1.

vim pod-pv.yml

apiVersion: v1

kind: persistentVolume

metadata:

name: my-pv

Spec:

StorageClassName: local-storage

Capacity:

Storage: 2Gi

accessModes:

- ReadWriteOnce.

hostPath:

Path: "/mnt/data"

→ kubectl create -f pod-pv.yml

kubectl get pods

kubectl get pv

* With the help of PV we need to create PVC,

vim pod-pvc.yaml

Step2:-

apiVersion: v1

kind: PersistentVolumeClaim

Metadata:

name: my-pvc

Spec:

storageClassName: local-storage

accessModes:

- ReadWriteOnce

resources:

requests:

Storage: 1Gi

kubectl create -f pod-pvc.yaml

kubectl get pods

kubectl get pvc

* Create a pod :-

Step3:- vim pod-hostpath.yaml

apiVersion: v1

kind: Pod

Metadata:

Name: pvc-pod

Spec:

Containers:

- name: busybox

image: busybox.

Command: ["/bin/sh"; "-c"; "while true;

do sleep 3600; done"]

VolumeMounts:

- mountPath: "/mnt/storage"

name: my-storage.

mount to this

location.

volumes:

- name: my-storage

persistentVolumeClaim:

claimName: my-pvc

: 1GiB

claimName: my-pvc : 1GiB

→ kubectl create -f pod-hostpath.yaml

kubectl get pvc.

kubectl get pods.

kubectl describe pods pvc-pod

Step 4:-
connect to a pod

kubectl exec -it pvc-pod -- bash

cd /mnt/storage.

*

Delete

kubectl delete pod pvc-pod.

kubectl delete pod my-pv.

Pod



* check Storage on Node1

master.
Node1
Node2

APP → C1

C2

(or) Node 2.

Not port IP address we can access

9/11/23:

K8's ConfigMaps & Secrets: [Dynamic running configuration]

- A configMap is an API object used to store non-confidential data in key-value pairs.
- It can bind configuration files, command-line arguments, environment variables, port numbers and each other config artifacts to your pod's container and system components at runtime (dynamically) at the pod runtime.
- These are useful for storing and sharing non-sensitive unencrypted configuration information.

Create a configmap:- vim configmap.yaml

```
↓  
[data]  
kind: ConfigMap  
apiVersion: v1
```

```
kind: ConfigMap  
apiVersion: v1
```

```
metadata:  
  name: my-configmap
```

```
  data:  
    mykey: myvalue
```

```
    anotherkey: anothervalue
```

→ kubectl create -f configmap.yaml

```
kubectl get cm
```

* Create a pod using cm:-

Vim configmap-pod.yaml

```
apiVersion: v1
kind: pod
metadata:
  name: my-configmap-pod
spec:
  containers:
    - name: my-container
      image: alpine
      command: ['sh', '-c', "echo ${my-var} && sleep 3600"]
  env:
    - name: my-var
      valueFrom:
        ConfigMapKeyRef:
          name: my-configmap
          key: mykey
```

firstly: kubectl create -f configmap-pod.yaml
kubectl get pods.

gradually: kubectl logs my-configmap-pod.

SECRETS:

- A Secret is an object that contains a small amount of sensitive data such as password, a token or SSH keys, authentication tokens etc.
- Secrets are similar to Configmaps but are specifically intended to hold confidential data.

→ User can create secrets and the system also creates some secrets.

User for secrets:-

- As files in a volume mounted on one or more of its containers.
- As container environment variables.
- By the kubectl

* managing a secret:- Several options to create a secret:

1. Using kubectl
2. Use a configuration file.
3. Use the kustomize tool

* Create a Secret Using kubectl Command :-

Step1: Encode a password

echo -n 'admin123' | base64

O/P: yWRtaWxzc2E=

Step2: Create a secret

→ kubectl create secret generic db-pass

--from-literal=username=admin

--from-literal=password=yWRtaWxzc2E=

→ kubectl get secrets

(type) http
opaque = encrypted

Decode a password :-

```
echo -n 'yWRtawuxmJm=' | base64 --decode
```

* Create a Secret using files :-

Step1. Create 2 files, with user and password

```
echo -n 'admin' > ./username.txt
```

```
echo -n 'yWRtawuxmJm=' > ./password.txt
```

```
ls. → cat username.txt
```

```
cat password.txt
```

Step2 :- Create a secret from 2 files :-

Kubectl Create Secret generic db-user-pass

```
--from-file=username=./username.txt
```

```
--from-file=password=./password.txt
```

→ kubectl get secrets.

→ To verify the Secrets: kubectl describe db-user-pass.

* To modify (or) Edit Secrets:

kubectl edit Secrets db-user-pass

* Create a pod :-

```
vim secrets-pod.yml
```

apiVersion: v1

kind: Pod

metadata:

name: myPod-volumeSecret

Spec:

containers:

- name: volumeSecret
image:忙的

image: redis

volumeMounts:

- name: foo

mountPath: "/opt/secrets"

readOnly: true

volumes:

- name: foo

secret:

Secret Name: mySecret-manifest

→ kubectl create -f Secrets-pod.yaml

→ kubectl get pods.

* To describe pod:

kubectl describe pod myPod-volumeSecret.

* Connect to a pod :-

kubectl exec -it myPod-volumeSecret -- bash

cd /opt/secrets

ls

Certifications :-

CloudKt

Certified Kubernetes Administrator.

11

Applications developer.

11

Security specialist,

19/01/23

→ Ansible is a open-source configuration management tool.

ANSIBLE

→ Ansible is a open-source software provisioning, configuration management and application deployment tool.

→ Ansible is simple to use for system administration & developers because it is built on python.

→ It is included in as a part of Federal distribution of linux (Redhat, centos, debian, CBL) in etc directory.

→ Ansible's architecture is "Agentless". Work is pushed to remote host.

When ansible executes

Ansible Automation

→ It is used by thousands of organization globally to help them automate IT task, such as Configuration management, provisioning, workflow orchestration, application deployment and life-cycle management.